DATA 8

Spring 2022

Lecture 30

Linear Regression

Announcements

Correlation Coefficient

The Correlation Coefficient *r*

- Measures ***linear*** association
- Based on standard units
- $-1 \le r \le 1$
  - $r = 1$: scatter is perfect straight line sloping up
  - $r = -1$: scatter is perfect straight line sloping down
- $r = 0$: No linear association; *uncorrelated*

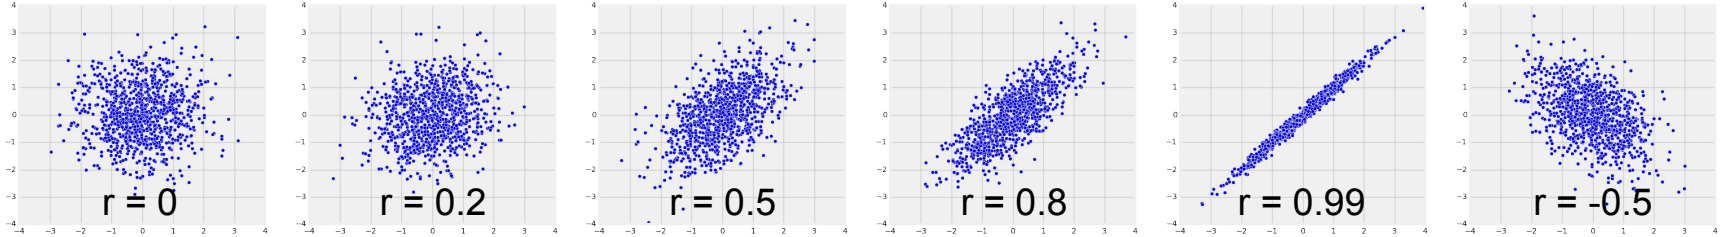

Definition of *r*

Correlation Coefficient (r) =

average of	product of	x in standard units	and	y in standard units

Measures how clustered the scatter is around a straight line

Care in Interpretation

Watch Out For ...

- False conclusions of causation
- Nonlinearity
- Outliers
- Ecological Correlations

(Demo)

Discussion question

True or False?

If the correlation of x and y is close to 0, then knowing one cannot help us predict the other.

Chocolate and Nobel Prizes

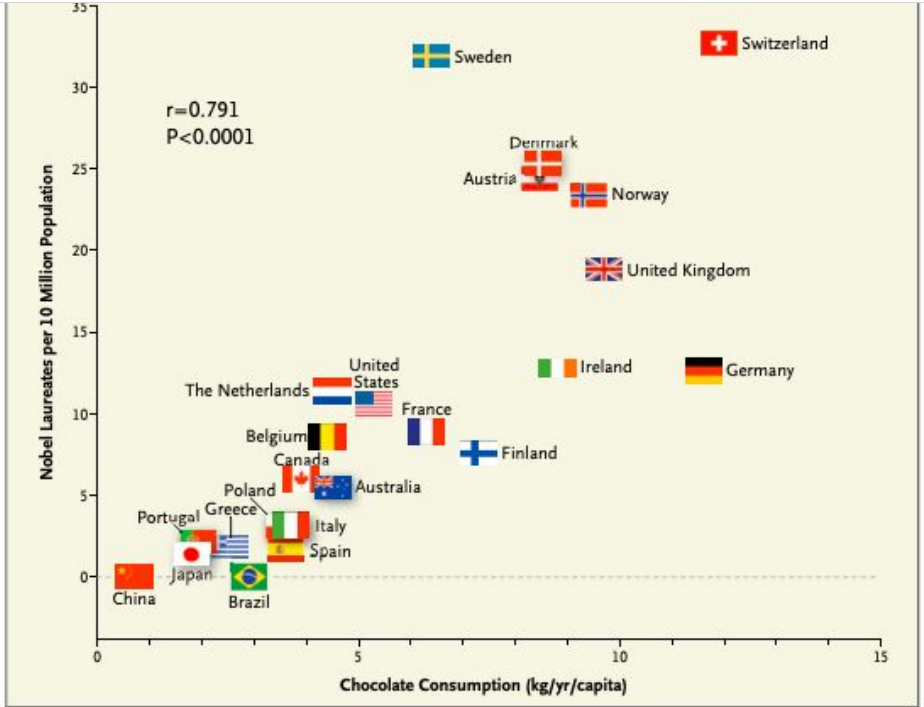

Figure 1. Correlation between Countries' Annual Per Capita Chocolate Consumption and the Number of Nobel Laureates per 10 Million Population.

https://www.biostat.jhsph.edu/courses/bio621/misc/Chocolate%20consumption%20cognitive%20function%20and%20nobel%20laurates%20(NEJM).pdf

Prediction

Predicting Heights

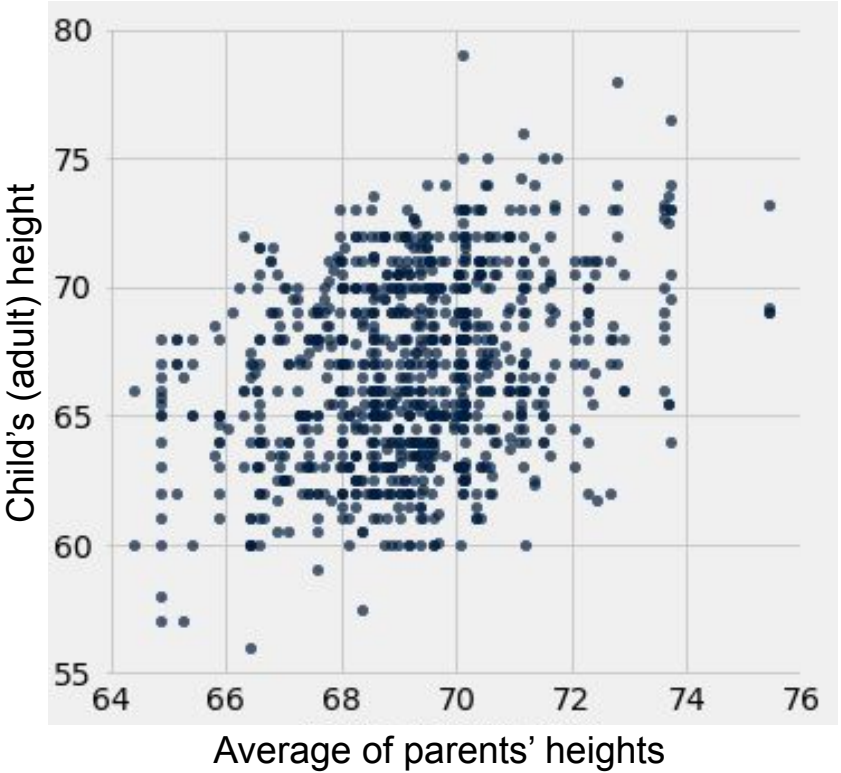

- Oval shaped
- Moderate positive correlation
- How can we predict child height from the parents' average height?

Approach to Prediction

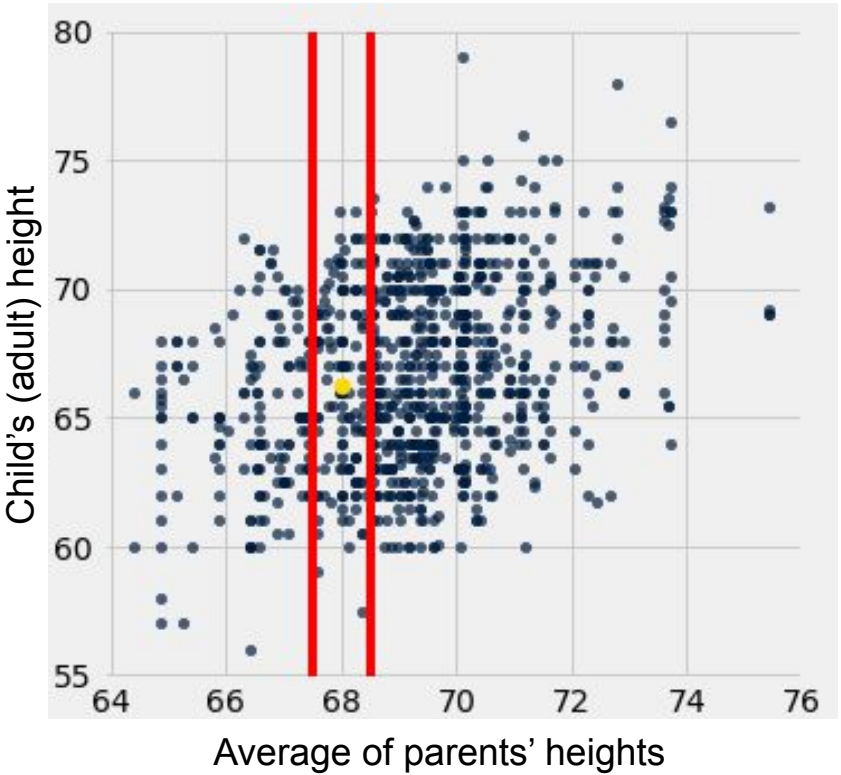

Predicted Heights

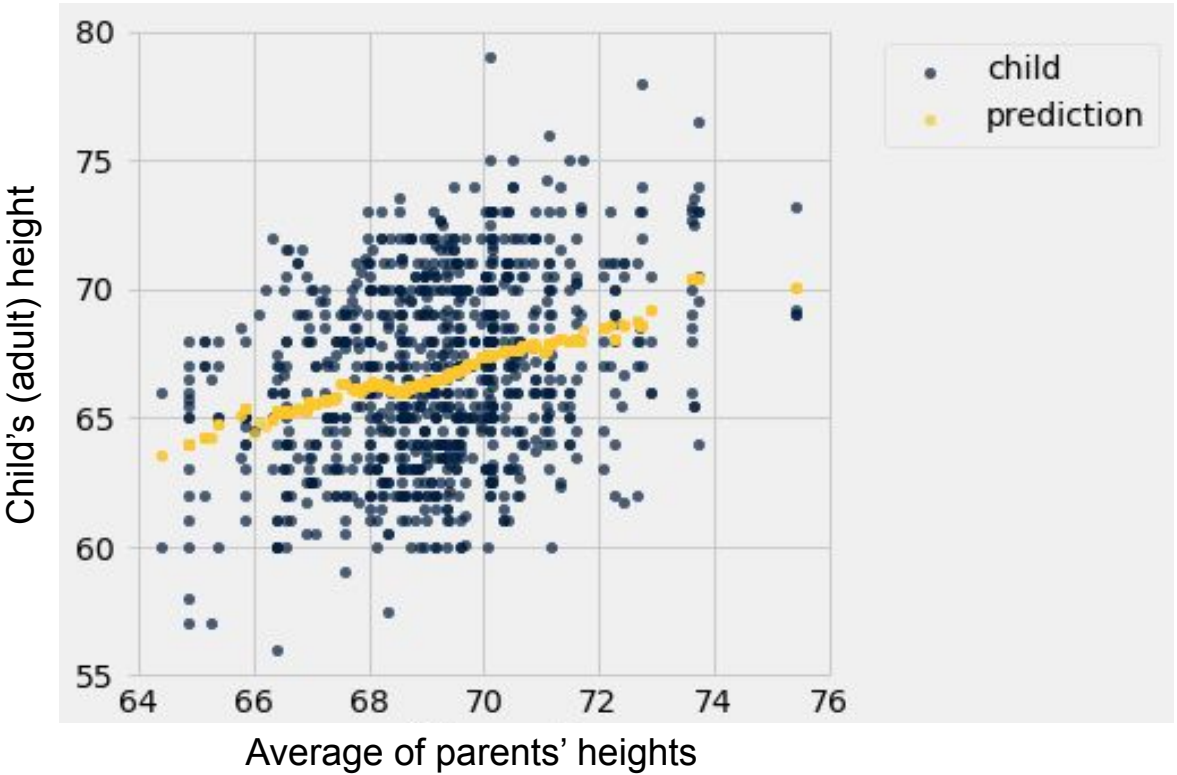

Nearest Neighbor Regression

A method for prediction:

- Group each *x* with similar (nearby) *x* values
- Average the corresponding *y* values for each group

For each *x* value, the prediction is the average of the *y* values in its nearby group.

The graph of these predictions is the “graph of averages”.

If the association between *x* and *y* is linear, then points in the graph of averages tend to fall on a line.

Where is the prediction line?

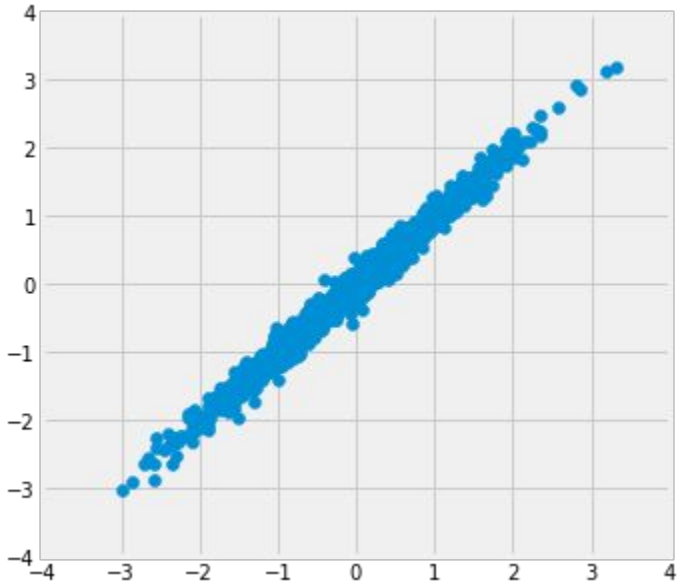

r = 0.99

Where is the prediction line?

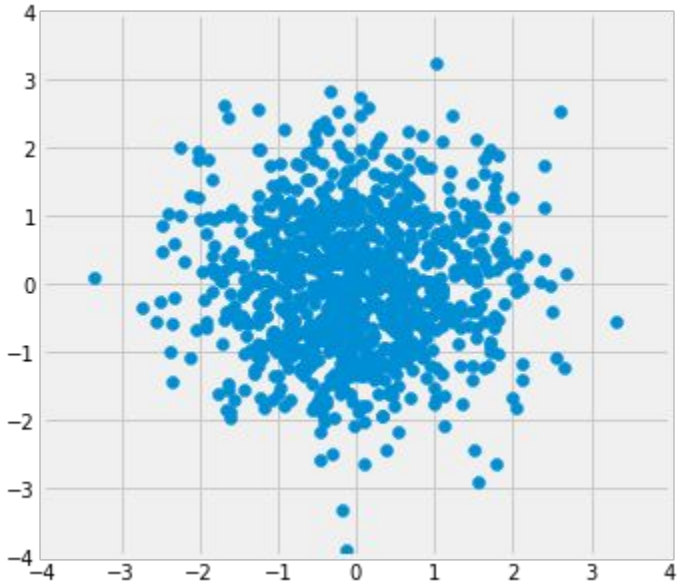

r = 0.0

(Demo)

Linear Regression

Linear Regression

A statement about *x* and *y* pairs

- **Measured in *standard units (su)***
- Describing the deviation of *x* from 0 (the average of *x*'s)
- And the deviation of the corresponding *y* from 0 (the average of *y*'s)

On average, *v* deviates from 0 less than *x* deviates from 0

Regression line

$$(\text{estimated } y)_{su} = r \times x_{su}$$

correlation

Not true for all points — a statement about averages